The Bigger, the Better? Optimizing Neural Networks for Calorimeter Calibration in the ATLAS Detector

Presenter: Annabel Li
Advisors: Colin Gay[1], Kelvin Leong[1,2], Wojtek Fedorko[2], Max Swiatlowski[2]

Background & Motivations

- Large Hadron Collider (LHC) will be upgraded to High-Luminosity by 2030
  - **Collisions per bunch crossing will increase**

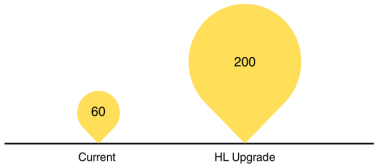

- Hardware trigger system (L0) already struggles with current data rate
  - Incorrect calibration in energy deposited → incorrect events reconstruction[1][2]
  - Low trigger rate discards potentially valuable information

We need a more accurate and efficient trigger system.

How Neural Networks Can Help

- DeepSets machine-learning model improves performance in cluster energy regression[1][2]
- 3 stages: Φ network, latent space, F network

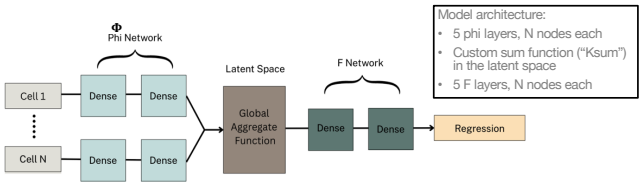

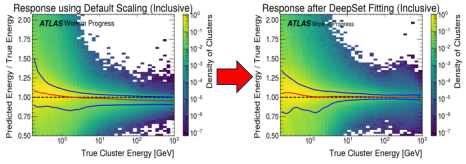

Fig1: DeepSets model visuals

Top: Schematic of the DeepSets model

Right: Response from MC samples using default calorimeter calibration (left) vs. DeepSets model (far right). Red/blue lines represent the median and IQR responses

How is Code Implemented on Hardware?

- FPGAs are designed with hardware description languages (VHDL, Verilog)
- **hls4ml** package automatically converts python machine learning models to synthesis-ready form

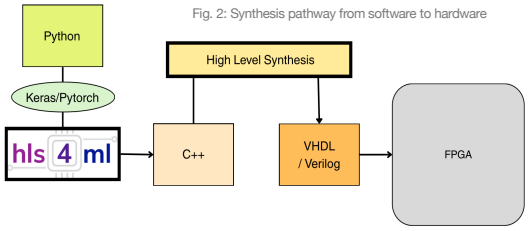

Fig. 2: Synthesis pathway from software to hardware

What is Quantization?

- During HLS, floating-point numbers are **quantized** to fixed
  - "ap_fixed<M,N>" = M total bits with N integer bits

ap_fixed<16,6> → **101101.1010000000 = -18.375**

- 2 methods for ML:
  - **Post-Training Quantization (PTQ)** → weights and biases quantized after training
  - **Quantization-Aware Training (QAT)** → model trained on lower-precision operations

We can use hls4ml to quickly test parameterizations of the DeepSets model for optimization[4][5].

References:

[1] "Deep Learning for Pion Identification and Energy Calibration with the ATLAS Detector," tech. rep., CERN, Geneva, 2020.
[2] "Point Cloud Deep Learning Methods for Pion Reconstruction in the ATLAS Experiment," tech. rep., CERN, Geneva, 2022.
[3] F. Fahim, et al., "hls4ml: An open-source codesign workflow to empower scientific low-power machine learning devices," March 2021.
[4] P. Odagiu, et al., "Ultrafast jet classification at the hl-lhc," Machine Learning: Science and Technology, vol. 5, p. 035017, July 2024.
[5] C. Antel, "QDIPS: Deep Sets Network for FPGA investigated for high-speed inference on ATLAS," tech. rep., CERN, Geneva, 2025.

Results

PTQ - varying number of integer bits

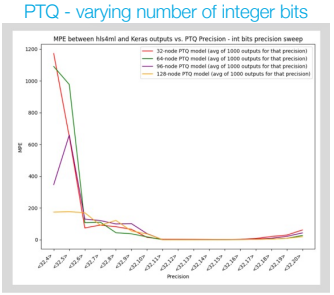

PTQ - varying number of total bits

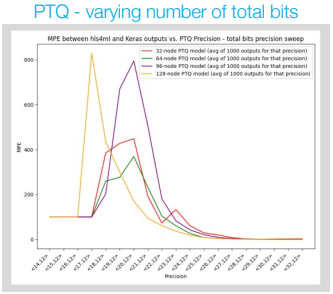

PTQ – setting architecture and intermediate output precisions separately

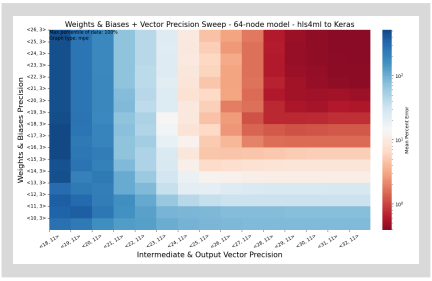

Fig3: PTQ results

Top left: MPE from Keras regression output, varying the number of integer bits for all model parameters

Top right: MPE from Keras regression output, varying the number of total bits for all model parameters

Left: MPE from Keras regression output, specifying different precisions for weights + biases and intermediate + output vectors

Conclusion & Next Steps

Problem: Current L0 trigger system at the LHC is unsuitable for the HL upgrade

Project Goal: Optimize NN size and precision for FPGA deployment

Findings:

- Larger models deviate more than smaller models from their Keras equivalents at lower precisions
- Accuracy increases with precision, but plateaus after a certain point
- Weights and biases can be represented with less bits than intermediate outputs

Next Steps:

- Further optimization strategies: QAT, pruning, High-Granularity Quantization